

Diabetes Analysis

Marina Huang

2023-04-21

The Data

In this project, I practice data analysis on health care data using R.

The dataset can be downloaded from [this link](#) (click). It is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of this dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. All patients here are females at least 21 years old of Pima Indian heritage.

This dataset consist of several medical predictor (independent) variables and one target (dependent) variable, **outcome**. Independent variables include the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

The summary of the dataset is displayed here:

```
summary(diabetes)
```

```
##      Pregnancies      Glucose      BloodPressure      SkinThickness
##  Min.   : 0.000    Min.   :  0.0    Min.   :  0.00    Min.   :  0.00
## 1st Qu.: 1.000    1st Qu.: 99.0    1st Qu.: 62.00    1st Qu.:  0.00
## Median : 3.000    Median :117.0    Median : 72.00    Median :23.00
## Mean   : 3.845    Mean   :120.9    Mean   : 69.11    Mean   :20.54
## 3rd Qu.: 6.000    3rd Qu.:140.2    3rd Qu.: 80.00    3rd Qu.:32.00
## Max.   :17.000    Max.   :199.0    Max.   :122.00    Max.   :99.00
##      Insulin      BMI      DiabetesPedigreeFunction      Age
##  Min.   :  0.0    Min.   :  0.00    Min.   :0.0780    Min.   :21.00
## 1st Qu.:  0.0    1st Qu.:27.30    1st Qu.:0.2437    1st Qu.:24.00
## Median : 30.5    Median :32.00    Median :0.3725    Median :29.00
## Mean   : 79.8    Mean   :31.99    Mean   :0.4719    Mean   :33.24
## 3rd Qu.:127.2    3rd Qu.:36.60    3rd Qu.:0.6262    3rd Qu.:41.00
## Max.   :846.0    Max.   :67.10    Max.   :2.4200    Max.   :81.00
##      Outcome
##  Min.   :0.000
## 1st Qu.:0.000
## Median :0.000
## Mean   :0.349
## 3rd Qu.:1.000
## Max.   :1.000
```

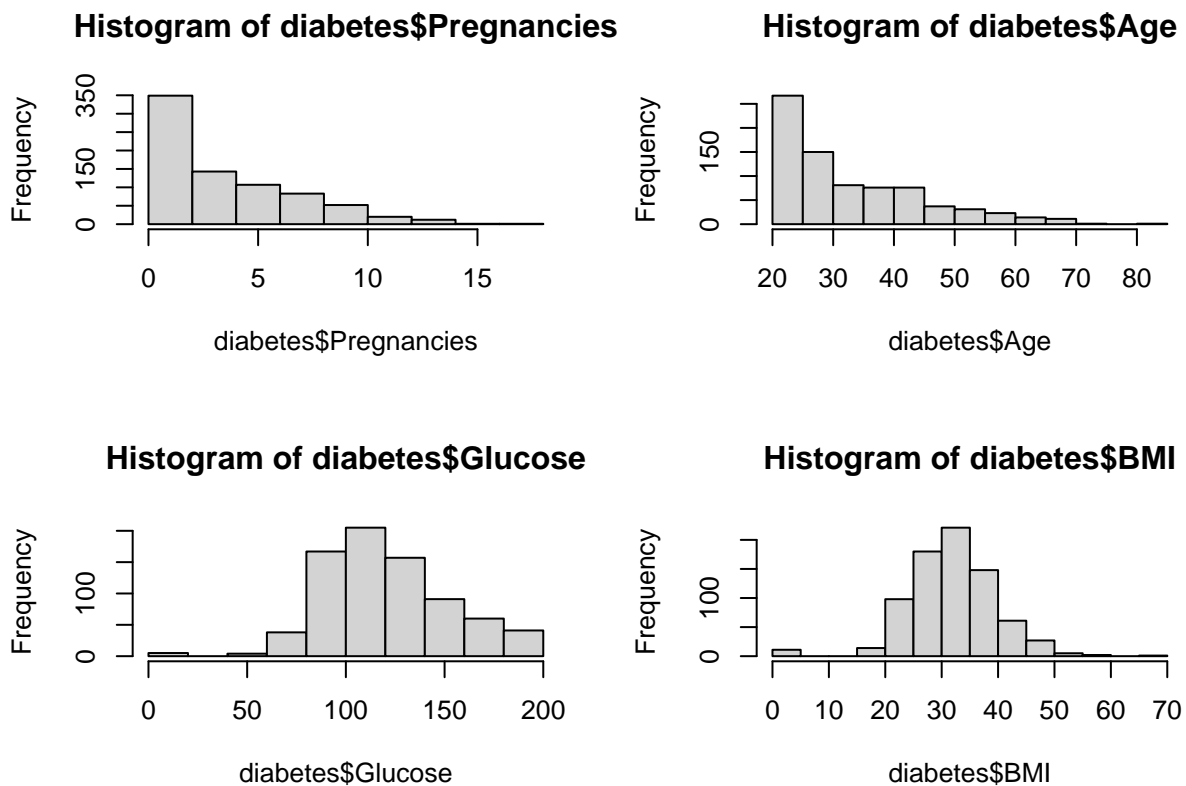
Explorative Data Analysis

Univariate Data Analysis

Univariate analysis explores each variable in a data set, separately. It looks at the range of values, as well as the central tendency of the values.

Let's take a look at the effect of **Pregnancies**, **Age**, **Glucose**, and **BMI** on diabetes:


```
par(mfrow = c(2,2))
hist(diabetes$Pregnancies)
hist(diabetes$Age)
hist(diabetes$Glucose)
hist(diabetes$BMI)
```



From these distribution graphs, **Age** and **Pregnancies** are not in normal distribution as expected, since the underlying population should not be normally distributed either. Glucose level and BMI are following a normal distribution.


Now looking into BloodPressure:

```
boxplot(diabetes$BloodPressure,  
        ylab = "Blood Pressure")
```



The impact of glucose on diabetes:

```
diabetes %>%  
  ggplot(aes(x = Glucose)) +  
  geom_histogram(fill = "sky blue", color = "black") +  
  facet_grid(Outcome~.)
```



Goal: Assess the mean difference of glucose levels between the positive and negative groups.

Null Hypothesis: There is no significant difference between glucose levels in positive and negative groups.

Conditions

- Individuals are independent of each other.
- Distribution is skewed (not normal), but there is >30 samples.
- Both the groups are independent of each other and the sample size is lesser than 10% of the population.

```
t.test(Glucose ~ Outcome, diabetes)
```

```
##
## Welch Two Sample t-test
##
## data: Glucose by Outcome
## t = -13.752, df = 461.33, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -35.74707 -26.80786
## sample estimates:
## mean in group 0 mean in group 1
##      109.9800      141.2575
```

The p-value is < 0.05 (the critical value), so we reject the null hypothesis for the alternate hypothesis. We can say that we are, 95% confident, that the average glucose levels for individuals with diabetes is $>$ the people without diabetes.

Box Plot of the impact of Age on Diabetes Pedigree Function

```
par(mfrow = c(1,2))

# Boxplot

with_d <- diabetes[diabetes$Outcome == 1,]
without <- diabetes[diabetes$Outcome == 0,]

boxplot(diabetes$DiabetesPedigreeFunction ~ diabetes$Outcome,
        ylab = "Diabetes Pedigree Function (DPF)",
        xlab = "Diabetes Presence",
        main = "Plot 1",
        outline = TRUE)

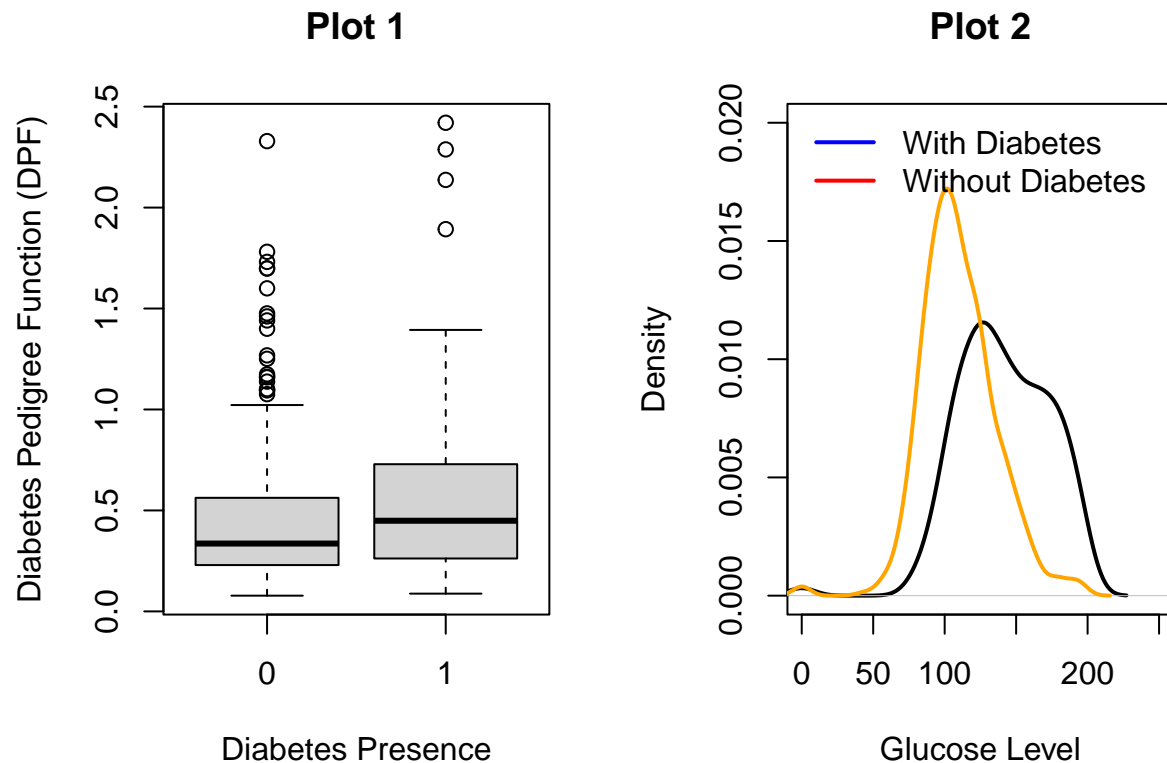
# Density Plot

plot(density(with_d$Glucose),
     xlim = c(0, 250),
     ylim = c(0.00, 0.02),
     xlab = "Glucose Level",
     main = "Plot 2",
     lwd = 2)

lines(density(without$Glucose),
     col = "orange",
     lwd = 2)

legend("topleft",
```

```
col = c("blue", "red"),
legend = c("With Diabetes", "Without Diabetes"),
lwd = 2,
bty = "n")
```



From Plot 2, the distribution is shifted towards the left for those without diabetes. This indicates those **without diabetes generally have a lower blood glucose level**.

Welch Two Sample t-Test

```
t.test(with_d$DiabetesPedigreeFunction, without$DiabetesPedigreeFunction)
```

```
##
## Welch Two Sample t-test
##
## data: with_d$DiabetesPedigreeFunction and without$DiabetesPedigreeFunction
## t = 4.5768, df = 454.51, p-value = 6.1e-06
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 0.06891135 0.17262065
## sample estimates:
## mean of x mean of y
## 0.550500 0.429734
```

Data Correlation Analysis

Scatter Matrix of All Columns:

```
ggcorr(diabetes[, -9], name = "Correlation", label = TRUE) +  
  theme(legend.position = "right") +  
  labs(title = "Correlation Plot of Variance") +  
  theme(plot.title = element_text(face = "bold", color = "black", hjust = 0.5, size = 12))
```

Correlation Plot of Variance



Pregnancy, Age, Insulin, SkinThickness are having higher correlation.

Conclusion

To understand what variables have an impact on diabetes, we first took a look at each variable separately (univariate data analysis). By displaying histograms of some dependent variables, we see that **Age** and **Pregnancies** are not normally distributed, which was to be expected since the underlying population should not be normally distributed by either. Glucose and BMI are following a normal distribution.

We conducted a Welch's Two Sample t-Test to determine if glucose levels between positive and negative groups of having diabetes had a significant difference. The p-value of this test was below the critical value of < 0.05 , so we rejected the null hypothesis. This means that we are 95% confident that the average glucose levels for individuals with diabetes is greater than the people without diabetes.

Knowing the results from the t-Test, we plotted a line graph of glucose levels in people with and without diabetes. The plot indicated that those without diabetes generally have a lower blood glucose level.

We can also look at the overall correlation between the variables with a correlation plot. Age, Pregnancies, Insulin, and Skin Thickness have high correlation.